

Understanding Artificial Neural Networks

Theory, Logic, and Python Implementation

AI Education Series

December 24, 2025

Lecture Overview

- 1 Introduction
- 2 Architecture
- 3 Activation Functions
- 4 Mathematical Logic
- 5 The Learning Process
- 6 Practical Implementation
- 7 Summary & Future

What is an Artificial Neural Network (ANN)?

- **Definition:** A computational model inspired by the biological neural networks of the human brain.
- **Core Purpose:** To simulate the brain's ability to learn and identify patterns from data.
- **Mechanism:** Consists of interconnected nodes (neurons) that process information using mathematical functions.

Biological vs. Artificial Neurons

Biological Neuron:

- *Dendrites*: Receive signals.
- *Soma*: Processes signals.
- *Axon*: Transmits output.

Human Neuron Diagram
Showing Dendrites, Cell Body
(Soma), and Axon.

Artificial Neuron:

- *Inputs*: Data points (x_n).
- *Weights/Bias*: Signal strength.
- *Activation*: Decision making.





Figure: Biological Structure

The Architecture: Layers

An ANN is typically organized into three types of layers:

- ① **Input Layer:** Receives raw data from the external world.
- ② **Hidden Layers:** Where the "learning" happens. These layers extract features.
- ③ **Output Layer:** Provides the final prediction or classification.

Visualizing the Structure




Activation: Sigmoid Function

Equation:

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

- **Range:** $(0, 1)$
- **Use Case:** Output layer for binary classification.
- **Issue:** Vanishing Gradient problem for large/small values.




Activation: Hyperbolic Tangent (Tanh)

Equation:

$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

- **Range:** $(-1, 1)$
- **Key Feature:** Zero-centered (outputs have mean near 0).
- **Use Case:** Often preferred over Sigmoid in hidden layers.




Activation: Rectified Linear Unit (ReLU)

Equation:

$$f(x) = \max(0, x)$$

- **Range:** $[0, \infty)$
- **Benefits:** Computationally efficient; reduces likelihood of vanishing gradient.
- **Standard:** Most common for Hidden Layers.




Activation: Leaky ReLU

Equation:

$$f(x) = \begin{cases} x & \text{if } x > 0 \\ 0.01x & \text{if } x \leq 0 \end{cases}$$

- **Purpose:** Fixes the "Dying ReLU" problem by allowing a small gradient when $x < 0$.
- **Range:** $(-\infty, \infty)$



The Mathematics of a Single Neuron

The output Y of a neuron is calculated as:

$$Y = f \left(\sum_{i=1}^n w_i x_i + b \right)$$

Where:

- x_i : Input signals.
- w_i : Synaptic weights (Importance of input).
- b : Bias (Threshold adjustment).
- f : Activation function (Sigmoid, ReLU, etc.).

Weights and Bias: The Learning Parameters

- **Weights (w):** Represent the strength of the connection. Higher weight means more influence.
- **Bias (b):** Shifts the activation function to allow flexibility in the model.
- **Goal of Training:** Finding optimal w and b to minimize error.

The Forward Pass (Inference)

- ① Data enters the input layer.
- ② Multiplied by weights, bias added.
- ③ Passed through activation function.
- ④ Moves to the next layer until it reaches the output.
- ⑤ **Result:** A raw prediction.

Measuring Error: Loss Functions

- **MSE (Mean Squared Error):** For regression.
- **Cross-Entropy:** For classification.

Backpropagation: The Chain Rule

- Propagates the error backward through the network.
- Calculates the gradient of the loss function with respect to weights.
- This is the core engine of AI learning.

Optimization: Gradient Descent

- Algorithm to update weights: $w_{new} = w_{old} - \alpha \frac{\partial Loss}{\partial w}$.
- **Learning Rate (α)**: Step size of the update.

Python Tools for Neural Networks

- **NumPy:** Fundamental package for scientific computing.
- **TensorFlow/Keras:** High-level API for building and training models.
- **PyTorch:** Popular research-focused deep learning library.

Step 1: Importing Libraries & Data

```
1 import tensorflow as tf
2 from tensorflow.keras.models import Sequential
3 from tensorflow.keras.layers import Dense
4 import numpy as np
5
6 # Sample Data (XOR Problem)
7 X = np.array([[0,0], [0,1], [1,0], [1,1]], "float32")
8 y = np.array([[0], [1], [1], [0]], "float32")
```

Step 2: Defining the Model

```
1 model = Sequential()  
2  
3 # Input layer + Hidden Layer (4 neurons)  
4 model.add(Dense(4, input_dim=2, activation='relu'))  
5  
6 # Output Layer (1 neuron for binary outcome)  
7 model.add(Dense(1, activation='sigmoid'))
```

Step 3: Compiling

```
1 model.compile(  
2     loss='binary_crossentropy',  
3     optimizer='adam',  
4     metrics=['accuracy'])  
5 )
```

Step 4: Training the Model (Fit)

```
1 # Training for 500 iterations (epochs)
2 history = model.fit(X, y, epochs=500, verbose=0)
3
4 # Evaluate performance
5 scores = model.evaluate(X, y)
6 print(f"Accuracy: {scores[1]*100}%")
```

Key Takeaways

- ① ANN is a bio-inspired mathematical model.
- ② Forward + Backpropagation is the learning cycle.
- ③ Python (Keras/TensorFlow) makes building networks accessible.
- ④ Preprocessing and tuning are as important as the architecture.

Thank you! Time for Coding!